

ADMFirstChallenge

```
#Importing Dataset
mnist <- read_mnist("~/Mscs 341 S22/Class/Data")
str(mnist)

## List of 2
## $ train:List of 2
## ..$ images: int [1:60000, 1:784] 0 0 0 0 0 0 0 0 0 0 ...
## ..$ labels: int [1:60000] 5 0 4 1 9 2 1 3 1 4 ...
## $ test :List of 2
## ..$ images: int [1:10000, 1:784] 0 0 0 0 0 0 0 0 0 0 ...
## ..$ labels: int [1:10000] 7 2 1 0 4 1 4 9 5 9 ...
```

Dataset Creation

- Your dataset should have in total 1000 randomly selected digits (feel free to use a set.seed command so that your results are reproducible).

First, calculate values for an individual image:

```
set.seed(12345) #For Reproducible Values

index <- sample(1:60000, 60000) #vector of randomly selected indexes to sample from mnist

tester <-
  as_tibble(index) %>%
  mutate(image = mnist$train$images[value, ]) %>%
  mutate(label = mnist$train$labels[value]) %>%
  dplyr::filter(label == 0 | label == 7)

tester$image[10,] #this pulls out a single image. This will be plotted to confirm if it's correct
```

```
## [1] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
## [19] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
## [37] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
## [55] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
## [73] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
## [91] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
## [109] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
## [127] 0 0 86 255 226 170 29 0 0 0 0 0 0 0 0 0 0 0
## [145] 0 0 0 0 0 0 0 0 0 0 0 0 86 255 170 170 255 226 114
## [163] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
## [181] 0 0 0 255 226 0 0 0 170 255 57 0 0 0 0 0 0 0 0
## [199] 0 0 0 0 0 0 0 0 0 114 255 198 29 0 255 57 0 0 0
## [217] 0 255 170 0 0 0 0 0 0 0 0 0 0 0 0 0 0 57 198
## [235] 255 198 255 226 0 255 86 0 0 0 0 170 226 0 0 0 0 0
## [253] 0 0 0 0 0 0 0 0 0 255 255 86 0 198 255 29 86 170 0
## [271] 0 0 0 141 255 29 0 0 0 0 0 0 0 0 0 0 0 0 29
```

```
## [289] 255 226 0 0 57 255 86 0 0 0 0 0 0 29 255 86 0 0
## [307] 0 0 0 0 0 0 0 0 0 0 170 255 114 0 0 57 255 86 0
## [325] 0 0 0 0 0 0 255 86 0 0 0 0 0 0 0 0 0 0 0
## [343] 0 226 255 57 0 0 141 170 0 0 0 0 0 0 0 0 29 255 86
## [361] 0 0 0 0 0 0 0 0 0 0 0 57 255 198 0 0 0 226 57
## [379] 0 0 0 0 0 0 0 0 86 255 0 0 0 0 0 0 0 0 0
## [397] 0 0 170 255 114 0 0 0 0 0 0 0 0 0 0 0 0 0 141
## [415] 255 0 0 0 0 0 0 0 0 0 0 0 0 170 255 86 0 0 0
## [433] 0 0 0 0 0 0 0 0 0 0 170 255 0 0 0 0 0 0 0
## [451] 0 0 0 0 170 255 86 0 0 0 0 0 0 0 0 0 0 0 0
## [469] 0 226 141 0 0 0 0 0 0 0 0 0 0 0 0 170 255 86 0
## [487] 0 0 0 0 0 0 0 0 0 0 0 86 255 86 0 0 0 0 0
## [505] 0 0 0 0 0 0 114 255 141 0 0 0 0 0 0 0 0 0 0
## [523] 0 0 198 226 0 0 0 0 0 0 0 0 0 0 0 0 0 0 255
## [541] 255 57 0 0 0 0 0 0 0 0 0 29 170 255 86 0 0 0
## [559] 0 0 0 0 0 0 0 0 0 0 170 255 226 29 0 0 0 0 0
## [577] 0 57 226 255 198 29 0 0 0 0 0 0 0 0 0 0 0 0 0
## [595] 0 29 226 255 198 114 86 86 141 226 255 255 255 114 0 0 0 0 0
## [613] 0 0 0 0 0 0 0 0 0 0 0 0 0 29 226 255 255 255
## [631] 255 255 198 114 29 0 0 0 0 0 0 0 0 0 0 0 0 0
## [649] 0 0 0 0 0 0 114 198 226 170 141 29 0 0 0 0 0 0
## [667] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
## [685] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
## [703] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
## [721] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
## [739] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
## [757] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
## [775] 0 0 0 0 0 0 0 0 0 0 0 0
```


```
tester$label[10] #Obtaining label to see if it matches with image. This is labelled 0.
```

```
## [1] 0
```

Plot the image to ensure data is extracted correctly. It should be a 0.

```
#Function that makes the image
plotImage <- function(dat,size=28){
  imag <- matrix(dat,nrow=size)[,28:1]
  image(imag,col=grey.colors(256), xlab = "", ylab="")
}

#Plotting the image of 0.
plotImage(tester$image[10,])
```



Looks like a 0. Data above was extracted correctly.

- Your training dataset should have 800 observations and your testing should have 200 observations.

Dividing the dataset into testing and training:

```
#Getting 1000 rows
tester <- tester %>%
  slice(1:1000)

#Setting Seed
set.seed(12345)

#Splitting into training and testing with respective ratio of 4:.1
trainer <- slice_sample(tester, n = 800)
tester <- setdiff(tester, trainer)
```

- Use the mnist dataset from dslabs to create an end-to-end classifier that distinguishes between 7 and 0.

Feature Definition

- You are allowed to use only 2 features. Notice that you need to calculate those features directly from dataset. Make sure to describe what those features represent and why you chose them. Are those features capturing any intuition that you have about distinguishing those two digits?

To determine the best way to calculate x_1 and x_2 , the quadrants will be calculated individually (q_1 , q_2 , q_3 , and q_4) so they can be used to calculate different possible combinations of x_1 and x_2 . We will test x_1 and x_2 based on which appears to work best when plotted on a scatter plot.

Features:

Feature 1: We test $x_1 = q_1 + q_3 - q_2 - q_4$ as a measure of symmetry. 0 should be much more symmetrical than 7, so values of x_1 should be smaller for 0.

Feature 2: x_2 is the sum of darkness in pixels of the top left quadrant q_2 . We expect this to be smaller for 7 than for 0 because the top-left area of a 7 is not very large.

First, the four quadrants will be calculated. Choose a single vector of 0: (will also do this process for 7 as a comparison: `tester$label[1]`)

```
zeroMatrix <- matrix(trainer$image[4, ], nrow=28)[,28:1]
sum(zeroMatrix[1:14, 15:28]) #q1
sum(zeroMatrix[1:14, 1:14])  #q2
sum(zeroMatrix[15:28, 1:14]) #q3
sum(zeroMatrix[15:28, 15:28]) #q4

#Observing the values of the quadrants+
abs(sum(zeroMatrix[1:14, 15:28]) + sum(zeroMatrix[15:28, 1:14]) - sum(zeroMatrix[1:14, 1:14]) - sum(zeroMatrix[15:28, 15:28]))

q1 = 6588;    q2 = 8456;    q3 = 7982;    q4 = 8562
```

Repeating this process above for 7:

```
sevenMatrix <- matrix(trainer$image[1, ], nrow=28)[,28:1]
sum(sevenMatrix[1:14, 15:28]) #q1
sum(sevenMatrix[1:14, 1:14])  #q2
sum(sevenMatrix[15:28, 1:14]) #q3
sum(sevenMatrix[15:28, 15:28]) #q4

abs(sum(sevenMatrix[1:14, 15:28]) + sum(sevenMatrix[15:28, 1:14]) - sum(sevenMatrix[1:14, 1:14]) - sum(sevenMatrix[15:28, 15:28]))

q1 = 4821;    q2 = 1757;    q3 = 7921;    q4 = 4487
```

The approach above ($q_1 + q_3 - q_2 - q_4$) gives 2448 for the zero and 6498 for the 7 which is about what we would expect - ideally the value returned for the zero image would be closer to zero. Similarly, the difference in the values of q_1 and q_2 is stark and apparent.

Next, calculate these values for all images and put it in a table:

Making the training dataset

```
#note that directly mutating these values into trainer or tester does not work. need to use a for-loop

trainVectorQ1 <- vector() #making sure the variables are clear
trainVectorQ2 <- vector()
trainVectorQ3 <- vector()
trainVectorQ4 <- vector()

for (i in 1:800) { #note: this value should be 1:200 for tester
  sumMatrix <- matrix(trainer$image[i, ], nrow = 28)[,28:1]
  q1 = sum(sumMatrix[1:14, 15:28]) #q1
```

```

print(q1)
trainVectorQ1 <- c(trainVectorQ1, q1)
trainVectorQ1
}
for (i in 1:800) {
  sumMatrix <- matrix(trainer$image[i, ], nrow = 28)[,28:1]
  q2 = sum(sumMatrix[1:14, 1:14])#q2
  trainVectorQ2 <- c(trainVectorQ2, q2)
  trainVectorQ2
}
for (i in 1:800) {
  sumMatrix <- matrix(trainer$image[i, ], nrow = 28)[,28:1]
  q3 = sum(sumMatrix[15:28, 1:14])#q3
  trainVectorQ3 <- c(trainVectorQ3, q3)
  trainVectorQ3
}
for (i in 1:800) {
  sumMatrix <- matrix(trainer$image[i, ], nrow = 28)[,28:1]
  q4 = sum(sumMatrix[15:28, 15:28])#q4
  trainVectorQ4 <- c(trainVectorQ4, q4)
  trainVectorQ4
}

```

```

trainer <- trainer %>%
  select(label) %>%
  mutate(label=as.factor(label)) %>%
  mutate(row = row_number()) %>%
  mutate(q1 = trainVectorQ1) %>%
  mutate(q2 = trainVectorQ2) %>%
  mutate(q3 = trainVectorQ3) %>%
  mutate(q4 = trainVectorQ4) %>%
  mutate(x_1 = abs(q1+q3-q2-q4)) %>%
  mutate(x_2 = q2)

```

trainer

```

## # A tibble: 800 x 8
##   label  row    q1    q2    q3    q4   x_1   x_2
##   <fct> <int> <int> <int> <int> <int> <int> <int>
## 1 7      1  4821  1757  7921  4487  6498  1757
## 2 7      2  6928   454  8455  4673 10256   454
## 3 0      3  2983  9026  5788  9287  9542  9026
## 4 0      4  6588  8456  7982  8562  2448  8456
## 5 0      5  2339  4741  3223  4547  3726  4741
## 6 0      6  9425 13349 10118 13699  7505 13349
## 7 7      7  1153  4153   893  4778  6885  4153
## 8 0      8  5438  9046  6584  8171  5195  9046
## 9 7      9  8020  3577  7285  8864  2864  3577
## 10 7     10  3584  7693  3360  8285  9034  7693
## # ... with 790 more rows

```


Plot the classifiers:

```

trainer %>%
  ggplot(aes(x = x_1, y = x_2, color = label)) +


```

```
geom_point()
```



x_1 and x_2 appear to successfully separate 0 and 7. There is not much overlap between the coordinate points of 0 and 7.

Repeat the process above to create the tester tibble and plot:



```
head(tester)
```

```
## # A tibble: 6 x 8
##   label row    q1    q2    q3    q4   x_1   x_2
##   <fct> <int> <int> <int> <int> <int> <int> <int>
## 1 7       1  2982  6035  2338  7285  8000  6035
## 2 7       2  7729  4013 10411  7379  6748  4013
## 3 0       3  5911  4780  6440  5859  1712  4780
## 4 7       4 12430   417 12408  9791 14630   417
## 5 0       5  7667  4910  8360  5266  5851  4910
## 6 7       6  6891  1977  7587  6635  5866  1977
```

This looks similar to the plot for trainer which is what is expected. Since there are fewer points, the difference between the points for 0 and 7 are not as defined. Next, these tables will be used to create and test models.

Model creation, optimization, and selection

- Create at least two different models for this classification and make sure to optimize the parameters those models have.
- Calculate the misclassification rates for both models and select the model with the lowest error rate.

Model #1

The first model we will be using is the logistic regression model from `parsnip` in `tidymodels`.

```
logit.model <- logistic_reg() %>%  
  set_engine("glm") %>%  
  set_mode("classification")  
  
default.recipe <-  
  recipe(label ~ x_1+x_2, data=trainer)  
  
logit.wflow <- workflow() %>%  
  add_recipe(default.recipe) %>%  
  add_model(logit.model)  
  
logit.fit <- fit(logit.wflow, trainer)  
logit.fit
```

Now, this model can be used on the tester dataset to classify whether images are 0 or 7 (in addition to the probability of being a value being 0 or 7).

```
predict(logit.fit, tester, type = "prob") #Gives probability of 0 and 7
```

```
## # A tibble: 200 x 2  
##       .pred_0 .pred_7  
##       <dbl> <dbl>  
##  1 0.0778    0.922  
##  2 0.0122    0.988  
##  3 0.166     0.834  
##  4 0.00000933 1.00  
##  5 0.0473    0.953  
##  6 0.00154    0.998  
##  7 0.00206    0.998  
##  8 0.884     0.116  
##  9 0.00223    0.998  
## 10 0.000208    1.00  
## # ... with 190 more rows
```

```
predict(logit.fit, tester) #Gives the classsification
```

```
## # A tibble: 200 x 1  
##       .pred_class  
##       <fct>  
##  1 7  
##  2 7  
##  3 7  
##  4 7  
##  5 7  
##  6 7  
##  7 7  
##  8 0  
##  9 7  
## 10 7  
## # ... with 190 more rows
```

Calculate the misclassification rate:


```
misclassification.tbl <- augment(logit.fit, tester)
mean(misclassification.tbl$label != misclassification.tbl$pred_class)
```

```
## [1] 0.115
```

The misclassification rate is 11.5%.

Model #2

Now, we are going to use K nearest neighbours (knn) from parsnip in tidymodels. `nearest_neighbor()` uses `k = 5` as default, let us attempt to optimize it

```
library(kknn)

#Making the model
knn.model <- nearest_neighbor(neighbors = tune()) %>%
  set_engine("kknn") %>%
  set_mode("classification")

#Making the workflow
knn.wflow <- workflow() %>%
  add_recipe(default.recipe) %>%
  add_model(knn.model)
```

Optimizing K:

```
#Making 10-fold cross-validation dataset
digits.folds <- vfold_cv(trainer, v = 10)
training(digits.folds$splits[2][[1]])
```

```
## # A tibble: 720 x 8
##   label row  q1  q2  q3  q4  x_1  x_2
##   <fct> <int> <int> <int> <int> <int> <int> <int>
## 1 7      1 4821 1757 7921 4487 6498 1757
## 2 7      2 6928 454 8455 4673 10256 454
## 3 0      4 6588 8456 7982 8562 2448 8456
## 4 7      7 1153 4153 893 4778 6885 4153
## 5 0      8 5438 9046 6584 8171 5195 9046
## 6 7      9 8020 3577 7285 8864 2864 3577
## 7 7     10 3584 7693 3360 8285 9034 7693
## 8 0     11 4188 7051 6351 8067 4579 7051
## 9 7     12 6906 5074 6678 8374 136 5074
## 10 7     13 4403 7409 5452 8317 5871 7409
## # ... with 710 more rows
```

```
testing(digits.folds$splits[2][[1]])
```


```
## # A tibble: 80 x 8
##   label row  q1  q2  q3  q4  x_1  x_2
##   <fct> <int> <int> <int> <int> <int> <int> <int>
## 1 0      3 2983 9026 5788 9287 9542 9026
## 2 0      5 2339 4741 3223 4547 3726 4741
## 3 0      6 9425 13349 10118 13699 7505 13349
## 4 0     21 10002 11298 9598 10906 2604 11298
## 5 0     31 8578 11081 7889 8720 3334 11081
## 6 7     32 7112 0 8203 5490 9825 0
```

```
## 7 0      33 4981 11160 7157 10279 9301 11160
## 8 7      36 2586 3558 5306 5712 1378 3558
## 9 0      48 6255 12299 8765 12700 9979 12299
## 10 7     52 4168 6451 3332 7120 6071 6451
## # ... with 70 more rows
```

```
#Making grid of neighbours across values of K
neighbors.tbl <- tibble(neighbors = seq(1,51, by = 1))
neighbors.grid.tbl <- grid_regular(neighbors(range = c(1, 51)),
                                  levels = 51)
```

```
#Tuning the results accordingly
tune.results <- tune_grid(object = knn.wflow,
                          resamples = digits.folds,
                          grid = neighbors.tbl)
```

```
#Having a look at the values of K
autoplot(tune.results)
```



```
#Show the best Value of K
show_best(tune.results, metric = "accuracy")
```

```
## # A tibble: 5 x 7
##   neighbors .metric .estimator mean      n std_err .config
##   <dbl> <chr>      <chr>    <dbl> <int>   <dbl> <fct>
## 1      14 accuracy binary    0.891    10 0.0110 Preprocessor1_Model14
## 2      13 accuracy binary    0.889    10 0.0116 Preprocessor1_Model13
```

```
## 3      15 accuracy binary    0.889    10 0.0106 Preprocessor1_Model15
## 4      11 accuracy binary    0.888    10 0.00986 Preprocessor1_Model11
## 5      12 accuracy binary    0.888    10 0.0113 Preprocessor1_Model12
```

```
best.neighbor <- select_best(tune.results, metric = "accuracy")
```

#Applying the optimal value of K (14)

```
knn.final.wflow <- finalize_workflow(knn.wflow, best.neighbor)
knn.fit <- fit(knn.final.wflow, trainer)
```

Finally, getting missclassification rate

```
predict(knn.fit, tester, type = "prob")
```

```
## # A tibble: 200 x 2
##   .pred_0 .pred_7
##   <dbl> <dbl>
## 1  0.0816 0.918
## 2  0.0867 0.913
## 3  0.332  0.668
## 4  0      1
## 5  0.0255 0.974
## 6  0      1
## 7  0      1
## 8  0.923  0.0765
## 9  0      1
## 10 0      1
## # ... with 190 more rows
```

```
predict(knn.fit, tester)
```

```
## # A tibble: 200 x 1
##   .pred_class
##   <fct>
## 1 7
## 2 7
## 3 7
## 4 7
## 5 7
## 6 7
## 7 7
## 8 0
## 9 7
## 10 7
## # ... with 190 more rows
```

```
misclassification.tbl <- augment(knn.fit, tester)
mean(misclassification.tbl$label != misclassification.tbl$.pred_class)
```

```
## [1] 0.125
```

Result of the models: Logistic Regression gave a misclassification rate of 11.5% and Knn gave 12.5%. Hence Logistic is marginally better. However logistic works with only 2 variables. Hence another model will be tested.

Model #3

```
library(tidymodels)
library(discrim)
tidymodels_prefer()

lda.model <- discrim_linear() %>%
  set_engine("MASS") %>% #MASS is the library, or a type of implementation
  set_mode("classification")

lda.wflow <- workflow() %>%
  add_recipe(default.recipe) %>%
  add_model(lda.model)

lda.fit <- fit(lda.wflow, trainer)
lda.fit

## == Workflow [trained] =====
## Preprocessor: Recipe
## Model: discrim_linear()
##
## -- Preprocessor -----
## 0 Recipe Steps
##
## -- Model -----
## Call:
## lda(..y ~ ., data = data)
##
## Prior probabilities of groups:
##      0      7
## 0.47375 0.52625
##
## Group means:
##      x_1      x_2
## 0 6084.715 10057.380
## 7 5187.404  4033.012
##
## Coefficients of linear discriminants:
##              LD1
## x_1  6.103812e-05
## x_2 -3.894835e-04

predict(lda.fit, tester, type="prob")

## # A tibble: 200 x 2
##   .pred_0 .pred_7
##   <dbl>   <dbl>
## 1 0.208   0.792
## 2 0.0489  0.951
## 3 0.171   0.829
## 4 0.000689 0.999
## 5 0.115   0.885
## 6 0.00937  0.991
## 7 0.0134   0.987
```

```
## 8 0.828      0.172
## 9 0.0107     0.989
## 10 0.00235   0.998
## # ... with 190 more rows
```

```
predict(lda.fit, tester)
```

```
## # A tibble: 200 x 1
##   .pred_class
##   <fct>
## 1 7
## 2 7
## 3 7
## 4 7
## 5 7
## 6 7
## 7 7
## 8 0
## 9 7
## 10 7
## # ... with 190 more rows
```

```
misclassification.tbl <- augment(lda.fit, tester)
mean(misclassification.tbl$label != misclassification.tbl$.pred_class)
```

```
## [1] 0.125
```

Misclassification = 12.5%

Visualization

- Plot the probabilities across a grid and the decision boundary for your selected model

As shown above, both Knn and LDA show an equally good misclassification rate. For the purpose of this project, we are going to use Knn. The probability of an image being 7 will be plotted with a decision boundary.

First, we found max and min values for `x_1` and `x_2` to create a grid for plotting with Bayes' boundary. Using max and min, the largest `x_1` is 20518, and the largest `x_2` is 18572. We create grids using seq up to 21000.


```
grid.vec.x_1 <- seq(from = 0, to = 21000, by = 500)
grid.vec.x_1
```

```
## [1]      0    500   1000   1500   2000   2500   3000   3500   4000   4500   5000   5500
## [13]  6000   6500   7000   7500   8000   8500   9000   9500  10000  10500  11000  11500
## [25] 12000  12500  13000  13500  14000  14500  15000  15500  16000  16500  17000  17500
## [37] 18000  18500  19000  19500  20000  20500  21000
```

```
grid.vec.x_2 <- seq(from = 0, to = 21000, by = 500)
grid.tbl <- expand_grid(x_1 = grid.vec.x_1, x_2 = grid.vec.x_2)
```

```
pred <- predict(knn.fit, grid.tbl, type="prob")#probability
pred %>%
  mutate(x_1 = grid.tbl$x_1) %>%
  mutate(x_2 = grid.tbl$x_2) %>%
  ggplot(aes(x_1, x_2, z=.pred_7, fill = .pred_7)) +
  geom_raster() +
```

```
stat_contour(breaks=c(0.5), color="pink")+
scale_fill_viridis_b()+
labs(title = "Decision Boundary Based on Knn Model",
fill = "Probability of 7")
```



Changing things up: adding the number 5

- Create a new dataset that includes your two chosen digits and the digit 5. Create training and testing datasets that include 5 and your two given digits.
- Calculate the same 2 features for this new testing and training dataset.


We copy and paste the code from near the beginning of this document and then plot it to confirm that everything looks as it should. This time, 5 is also included in the testing and training tables.

Training Dataset:

Having a look at the split between the parameters and the numbers

```
trainer %>%
  ggplot(aes(x = x_1, y = x_2, color = label)) +
  geom_point()+
  labs(title="Scatterplot of Features for Training Data With 0, 5, & 7")
```

Scatterplot of Features for Training Data With 0, 5, & 7



5 Appears to be right in the middle of the split of 0 and 7.

Testing Dataset:

```
#Making the model
knn.model <- nearest_neighbor(neighbors = tune()) %>%
  set_engine("knn") %>%
  set_mode("classification")

#Making the new Recipe
default.recipe <- recipe(label ~ x_1+x_2, data=trainer)

#Making the workflow
knn.wflow <- workflow() %>%
  add_recipe(default.recipe) %>%
  add_model(knn.model)
```

Optimizing K:

```
#Making 10-fold cross-validation dataset
digits.folds <- vfold_cv(trainer, v = 10)
training(digits.folds$splits[2][[1]])
```

```
## # A tibble: 720 x 8
##   label  row  q1    q2    q3    q4  x_1  x_2
##   <fct> <int> <int> <int> <int> <int> <int> <int>
## 1 5      1  5321  8600  3187  7124  7216  8600
## 2 5      2  9885  6952  9312 10396  1849  6952
## 3 5      4  7586  7314 11195  8544  2923  7314
```

```
## 4 5      7 5481 8984 8312 7936 3127 8984
## 5 7      8 11314 4358 12376 9729 9603 4358
## 6 7      9 7966    0 11823 7284 12505    0
## 7 7     10 7344    8 9854 5677 11513    8
## 8 5     11 9911 5725 10106 8094 6198 5725
## 9 5     12 9662 10191 11176 10790 143 10191
## 10 0     13 6611 12409 7096 12121 10823 12409
## # ... with 710 more rows
```


```
testing(digits.folds$splits[2][[1]])
```

```
## # A tibble: 80 x 8
##   label  row  q1    q2    q3    q4  x_1  x_2
##   <fct> <int> <int> <int> <int> <int> <int>
## 1 0      3 11956 15078 13652 15784 5254 15078
## 2 0      5 6134  9726  8635 11407 6364  9726
## 3 0      6 9180  8162  8124 10844 1702  8162
## 4 7     21 4522  3276  4380  6828 1202  3276
## 5 5     31 8774  5262  7823  5785 5550  5262
## 6 5     32 3104  4444  4526  4493 1307  4444
## 7 7     33 2376  1980  3842  4000  238  1980
## 8 0     36 5925 11427  8509 10144 7137 11427
## 9 7     48 4409  7759  3274 10162 10238 7759
## 10 0     52 5588 10029 7208  9249 6482 10029
## # ... with 70 more rows
```

```
#Making grid of neighbours across values of K
neighbors.tbl <- tibble(neighbors = seq(1,51, by = 1))
neighbors.grid.tbl <- grid_regular(neighbors(range = c(1, 51)),
                                   levels = 51)
```

```
#Tuning the results accordingly
tune.results <- tune_grid(object = knn.wflow,
                          resamples = digits.folds,
                          grid = neighbors.tbl)
```

```
#Having a look at the values of K
autoplot(tune.results)
```

#Show the best Value of K

```
show_best(tune.results, metric = "accuracy")
```

```
## # A tibble: 5 x 7
##   neighbors .metric .estimator mean      n std_err .config
##   <dbl> <chr>      <chr>      <dbl> <int>  <dbl> <fct>
## 1      49 accuracy multiclass 0.645    10  0.0155 Preprocessor1_Model49
## 2      48 accuracy multiclass 0.644    10  0.0157 Preprocessor1_Model48
## 3      50 accuracy multiclass 0.644    10  0.0153 Preprocessor1_Model50
## 4      51 accuracy multiclass 0.644    10  0.0153 Preprocessor1_Model51
## 5      45 accuracy multiclass 0.643    10  0.0148 Preprocessor1_Model45
```

```
best.neighbor <- select_best(tune.results, metric = "accuracy")
```

#Applying the optimal value of K (14)

```
knn.final.wflow <- finalize_workflow(knn.wflow, best.neighbor)
```

```
knn.fit <- fit(knn.final.wflow, trainer)
```

Calculating Misclassification rate:

```
augment(knn.fit, tester) %>%
  accuracy(truth = label, estimate = .pred_class)
```

```
## # A tibble: 1 x 3
##   .metric .estimator .estimate
##   <chr>    <chr>      <dbl>
## 1 accuracy multiclass    0.665
```

We get a misclassification of 33.5%

Confusion Matrix:

```
augment(knn.fit, tester) %>%
  conf_mat(truth = label, estimate = .pred_class)
```

```
##           Truth
## Prediction  0  5  7
##           0 56 13  5
##           5  9 29 20
##           7  3 17 48
```

5 appears to be problematic as a lot of them seem to be confused for 7s and a comparable amount seem to be confused with 0s. However, the model seems to work especially well for 0 and 7 isn't that bad either. 7 is very rarely confused for a 0 but is often confused for a 5.

Results:

Accuracy of 0: $56/(56+9+3) = 0.8235 \rightarrow 82.3\%$ Very good!

Accuracy of 5: $29/(13+29+17) = 0.4915 \rightarrow 49.1\%$ Mediocre

Accuracy of 7: $48/(5+20+48) = 0.6575 \rightarrow 65.7\%$ Fair

As we might expect, since our features are based on 0 and 7, the 5s are misclassified the most. Out of the misclassifications of 5s, $26/37 = 70\%$ were classified as 7s. The other 30% were misclassified as 0s. The reason why our model thinks 5s are 7s more often is because, like the average 7, the average 5 is not going to be perfectly symmetric, and we can expect most 5s to have some of their image in the upper left corner. If we ran our model with just 0s and 5s, we would expect to have a misclassification rate similar to our model for classifying 0s and 7s.

Creating the Grid:

```
create_grid <- function(delta) {
  expand_grid(x_1=seq(0,21000, by=delta), x_2 = seq(0,21000, by=delta))
}
grid.tbl <- create_grid(200)
grid.tbl <- grid.tbl %>%
  mutate(x_1 = as.integer(x_1)) %>%
  mutate(x_2 = as.integer(x_2))

augment.tbl <- predict(knn.fit, grid.tbl) %>%
  mutate(row = row_number())
augment2.tbl <- predict(knn.fit, grid.tbl, type = "prob") %>%
  mutate(row = row_number())

augment.tbl <- full_join(augment.tbl, augment2.tbl, by = "row") %>%
  mutate(x_1 = grid.tbl$x_1) %>%
  mutate(x_2 = grid.tbl$x_2)
augment.tbl
```

```
## # A tibble: 11,236 x 7
##   .pred_class row .pred_0 .pred_5 .pred_7 x_1 x_2
##   <fct>      <int>   <dbl>   <dbl>   <dbl> <int> <int>
## 1 7          1 0.00916 0.0183 0.973    0    0
## 2 7          2 0.0117 0.0208 0.968    0   200
## 3 7          3 0.0150 0.0208 0.964    0   400
## 4 7          4 0.0158 0.0225 0.962    0   600
```

```
## 5 7          5 0.0175  0.0233  0.959    0  800
## 6 7          6 0.0183  0.0250  0.957    0 1000
## 7 7          7 0.0200  0.0283  0.952    0 1200
## 8 7          8 0.0217  0.0292  0.949    0 1400
## 9 7          9 0.0233  0.0333  0.943    0 1600
## 10 7         10 0.0242  0.0350  0.941    0 1800
## # ... with 11,226 more rows
```

Plot Grid With Boundary

```
augment.tbl %>%
  ggplot() +
    geom_raster(aes(x_1, x_2, fill = .pred_class)) +
    geom_point(data=tester, aes(x=x_1, y=x_2, color=label, shape=label))+
    scale_color_manual(values=c("blue","red","orange"))+
    labs(title="Decision Boundary of LDA for 0, 5, & 7",
         fill = "Predicted Class",
         shape = "Label")
```

